

Bash Workshop I: The Basics

Frederick Yin

JITech

2023

Table of Contents

Intro

I. File Tree

II. CLI

III. Pipes

A brief history of bash



- ▶ Born: 1989
- ▶ Probably played Pokémon on the Game Boy
- ▶ Is an umbrella term for zsh, fish, ...
- ▶ Runs on Unix-like environments

A brief history of Unix

UNIX[®]

An Open Group Standard

- ▶ Born: 1969
- ▶ Probably listened to Michael Jackson
- ▶ Gave rise to Linux, BSD, and Mac OS
- ▶ We call them “Unix-like”

Unix: The Good Part

The Unix philosophy (paraphrased):

- ▶ Store data in plain text
- ▶ Hierarchical file system
- ▶ Everything is a file
- ▶ One tool does one thing
- ▶ Tools together strong

Quote

The power of a system comes more from the relationships among programs than from the programs themselves.

— Brian Kernighan and Rob Pike ¹

¹The UNIX Programming Environment. 1984. viii

Unix: The Chaotic Part

“Unix” is mostly created by these three groups of people who routinely disagree with each other:²

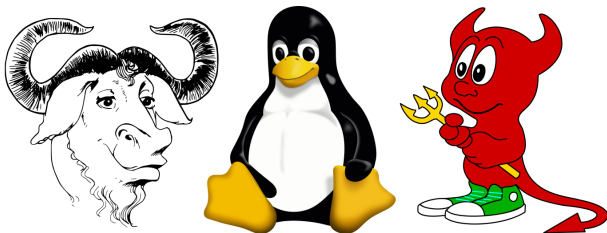


Figure: GNU, Linux, and BSD³

Despite this, they all hate Microsoft.
We will be talking about GNU today.

²Not convinced? Try `man ps`.

³Yes, I know there are many BSDs

Before we start

- ▶ This is **not** a Linux workshop (although I encourage you to use it)
- ▶ This is not a vim workshop either
- ▶ You should use a monospace font
- ▶ Anyone caught using PowerShell will be kicked out of the venue

Conventions in slides

- ▶ \$ indicates a **bash command**. Do not type the \$.
- ▶ # indicates a **comment**. Do not type it or anything after it.

For example, when you see:

```
1 $ echo hello # printf("hello\n");
```

You are going to type:

```
1 echo hello
```

Then hit Enter. I encourage you to type commands by hand.

Table of Contents

Intro

I. File Tree

II. CLI

III. Pipes

Files

Each of these is a different **file**:

- ▶ a
- ▶ .a (Hidden)
- ▶ a.txt
- ▶ A.txt
- ▶ A.TXT

Note

The dot and suffix are part of the filename.

Avoid spaces and special characters (except `._-`). If you have to, surround filename in quotes: 'Lab Report (3) final FINAL-1.docx'

cat: Printing a file

Open a bash terminal inside 01-files/, then:

```
1 $ cat a
2 $ cat .a
3 $ cat a.txt
```

Explanation

cat is short for “concatenate” (to join together) but it’s mostly used to print files.

cp, mv, rm: Relocating a file

Try this inside 01-files/:

```
1 $ ls
2 $ cp a b
3 $ ls
4 $ mv a.txt b.txt
5 $ ls
6 $ rm b
7 $ ls
```

Explanation

- ▶ `ls` lists files
- ▶ **copy** `a` into a file called `b`
- ▶ **move** `a.txt` into a file called `b.txt`
- ▶ **remove** `b`

cp, mv: Overwriting and renaming

When the destination does not exist, cp and mv simply create that file. **Otherwise, it is destroyed and overwritten.**

Try this inside 01-files/:

```
1 $ cp a b          # creates b
2 $ cat a
3 $ cp a a.txt      # overwrites a.txt
4 $ cat a.txt
```

Renaming a file in bash works like so:

```
1 $ mv b newb
```

Directories

Each of these is a **directory** (“dir” for short):

- ▶ 01-files/
- ▶ 01-files/c/
- ▶ 01-files/.c/ (Hidden dir)

Convention

For clarity, we add a slash (/) to the end of a directory in the slides. However, in reality it often makes no difference.

cd, pwd: Changing directory

Try this inside 01-files/:

```
1 $ cd c/  
2 $ pwd  
3 $ cd ../  
4 $ pwd
```

Explanation

- ▶ cd: “change directory”
- ▶ pwd: “print working directory”
- ▶ ../ means “parent directory”

ls: Listing directories

Try this inside 01-files/:

```
1 $ ls
2 $ ls -a
3 $ ls -l
4 $ ls -la
5 $ ls c/
```

Explanation

- ▶ ls: “list”
- ▶ -a is short for --all
- ▶ -l enables long listing format
- ▶ -la = -l + -a

mkdir, rm: Creating and deleting directories

Try this inside 01-files/:

```
1 $ mkdir dir/  
2 $ rm -r dir/  
3 $ mkdir -p dir/subdir/  
4 $ ls dir/
```

Explanation

- ▶ mkdir: “make directory”
- ▶ -r is short for --recursive
- ▶ -p is short for --parents

cp, mv: Into and out of directories

Try this inside 01-files/:

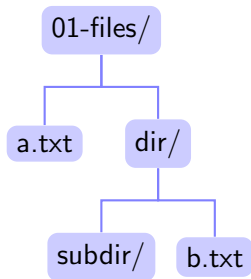
```
1 $ cp a dir/  
2 $ mv b.txt dir/  
3 $ mv dir/a b.txt
```

Explanation

- ▶ Copy a into dir/
- ▶ Move b.txt into dir/
- ▶ Move dir/a back into b.txt

File tree

Think of any directory as a tree.



Paths

File \cup directory = **path**.⁴

No paths under the same directory can bear the same name.

These **cannot** coexist:

- ▶ 01-files/data/, a directory
- ▶ 01-files/data, a regular file

⁴At least in the scope of this workshop.

Absolute & relative paths

- ▶ Paths beginning with `/` are absolute: `/usr/bin/cat`
- ▶ Otherwise it is relative: `01-files/`

If you know where you are, you can convert a relative path to an absolute one.

Example

Your location: `/home/you/`

Relative path: `bash-workshop/01-files/`

Absolute path: `/home/you/bash-workshop/01-files/`

Wildcard

* is a character to match any number of (including zero) characters.

Exception

Hidden paths will remain hidden unless you explicitly specify the dot: .*

Example

	a/	b/	a-copy.txt	b.txt
*	✓	✓	✓	✓
a*	✓		✓	
*.txt			✓	✓

Technically it's called a glob pattern but who cares. Also there are other weird symbols like ? or [] but I swear * is most of us will ever use.

. and ..

Inside every dir⁵ there are two special dirs:

- ▶ ./ — current dir
- ▶ ../ — parent dir

You can use them in relative paths.

Example

Your location: /home/you/

Relative path: ../friend/bash-workshop/01-files/

(Note that /home/you/../friend/ is just /home/friend/)

Absolute path: /home/friend/bash-workshop/01-files/

⁵Except /

Challenge

Inside 01-files/:

- ▶ Enter challenge/
- ▶ Create backup/
- ▶ Copy a.txt into dir/
- ▶ Move dir/ into backup/
- ▶ Verify using ls
- ▶ Delete backup/
- ▶ Go to the next section's directory

Solution

```
1 $ cd challenge/
2 $ mkdir backup/
3 $ cp a.txt dir/
4 $ mv dir/ backup/
5 $ ls
6 # Output: backup/  a.txt
7 $ ls backup/dir/
8 # Output: a.txt
9 $ rm -r backup/
10 $ cd ../../02-cli/
```

Table of Contents

Intro

I. File Tree

II. CLI

III. Pipes

The CLI

CLI stands for **command line interface**, as opposed to a GUI.

```
root@localhost ~# ping -q fo.wikipedia.org
PING text.ptpa.wikipedia.org (208.86.152.2) 56(84) bytes of data:
.
.
.
text.ptpa.wikipedia.org ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 540.528/540.528/540.528/0.000 ms
root@localhost ~# pwd
/root
root@localhost ~# cd /var
root@localhost var# ls -la
total 72
drwxr-xr-x. 18 root root 4096 Jul 30 22:43 .
drwxr-xr-x. 22 root root 4096 Sep 14 20:42 ..
drwxr-xr-x. 2 root root 4096 May 14 00:15 account
drwxr-xr-x. 11 root root 4096 Jul 31 22:26 cache
drwxr-xr-x. 3 root root 4096 May 18 16:03 db
drwxr-xr-x. 3 root root 4096 May 18 16:03 empty
drwxr-xr-x. 2 root root 4096 May 18 16:03 games
drwxr-x--T. 2 root gdm 4096 Jun 2 18:39 gdm
drwxr-xr-x. 38 root root 4096 May 18 16:03 lib
drwxr-xr-x. 2 root root 4096 May 18 16:03 local
drwxr-xr-x. 1 root root 11 May 14 00:12 lock -> ../run/lock
drwxr-xr-x. 14 root root 4096 Sep 14 20:42 log
drwxr-xr-x. 1 root root 10 Jul 30 22:43 mail -> spool/mail
drwxr-xr-x. 2 root root 4096 May 18 16:03 nis
drwxr-xr-x. 2 root root 4096 May 18 16:03 opt
drwxr-xr-x. 2 root root 4096 May 18 16:03 preserve
drwxr-xr-x. 2 root root 4096 Jul 1 22:11 report
drwxr-xr-x. 1 root root 6 May 14 00:12 run -> ../run
drwxr-xr-x. 14 root root 4096 May 18 16:03 spool
drwxr-xr-x. 4 root root 4096 Sep 12 23:56 tmp
drwxr-xr-x. 2 root root 4096 May 18 16:03 yp
root@localhost var# yum search wiki
Loaded plugins: langpacks, presto, refresh-packagekit, remove-with-leaves
python-free-updates | 2.7 kB | 00:00
python-free-updates/primary_db | 206 kB | 00:04
python-nonfree-updates | 2.7 kB | 00:00
updates/metalink | 5.9 kB | 00:00
updates | 4.7 kB | 00:00
updates/primary_db | 73% [=====] | 62 kB/s | 2.6 MB | 00:15 ETA
```

Figure: A stereotypical, Hollywood-like CLI.

Anatomy of a command

```
1 # get first 5 lines of file
2 $ head -n 5 longfile.txt
```

head	Executable file somewhere
-n	Option (aka flag)
5	Argument to -n
longfile.txt	Argument to head

Anatomy of a command

Anatomy of a command:

```
1 # list all files but *.o recursively in reverse
2 $ ls -Rr --ignore='*.o'
```

ls	Executable in \$PATH
-Rr	Two options: -R -r
--ignore=	Long option
'*.o'	Argument to --ignore

Notes

- ▶ Not every program uses this --long-option convention
- ▶ The equal sign after --ignore is optional in this command
- ▶ '*.o' does *not* expand to a list of files. It is simply a string.

I can't possibly remember all `--this` and `--that`!

You don't need to, thanks to **man pages**! (Short for manual pages)

Try:

```
1 $ man ls
```

If it doesn't work, try <https://man.archlinux.org/man/lis.1>

Challenge

- ▶ Read the man page for head
- ▶ Experiment with files in 02-cli/
- ▶ Find a command to generate the following:

```
1 ==> p0.txt <==  
2 MANIFESTO OF THE COMMUNIST PARTY.  
3  
4 ==> p1.txt <==  
5 I.  BOURGEOIS AND PROLETARIANS.
```

Challenge

- ▶ Read the man page for head
- ▶ Experiment with files in 02-cli/
- ▶ Find a command to generate the following:

```
1 ==> p0.txt <==  
2 MANIFESTO OF THE COMMUNIST PARTY.  
3  
4 ==> p1.txt <==  
5 I.  BOURGEOIS AND PROLETARIANS.
```

Solution

```
1 $ head -n1 -v p0.txt p1.txt
```


Learning by doing

Inside 02-cli/, run:

```
1 $ diff sway.1.conf sway.2.conf
```

Congratulations, you just learned how to use diff!

Now, what does this command do?

```
1 $ comm -12 sway.1.conf sway.2.conf
```

Lifehacks⁶

- ▶ Use ↑↓
- ▶ Ctrl-W: delete one word to the left
- ▶ Ctrl-U: delete everything to the left
- ▶ Ctrl-K: delete everything to the right
- ▶ Ctrl-7: undo (might not work in Git Bash)
- ▶ Ctrl-C: abort
- ▶ Ctrl-R: search history
- ▶ Ctrl-L: clear screen

⁶Should work in most shells.

Table of Contents

Intro

I. File Tree

II. CLI

III. Pipes

stdout

When you `printf`, where does the string go?

Your screen? Yes but also no.

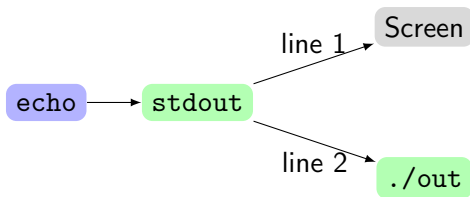
`stdout`, or **standard output**, is a special file where `printf` dumps its output.⁷

⁷Yes, there's `stderr` too, but we won't be talking about it.

Capturing stdout

You can capture the stdout of a command and direct it somewhere else than your screen, such as a file. Try this:

```
1 $ echo hello          # print to terminal
2 $ echo hello > out    # write to file
```



> and >>

Try >> instead of >. What happens?

```
1 $ echo hello >> out
2 $ cat out
```

Repeat a few times with both > and >>. What's the difference?

> and >>

Try >> instead of >. What happens?

```
1 $ echo hello >> out
2 $ cat out
```

Repeat a few times with both > and >>. What's the difference?

Observation

> overwrites the file, but >> appends to it.

stdin

The opposite of `stdout` is `stdin`: standard input. Some programs read from `stdin` when they expect a path but aren't given any.

Try this in `03-pipes/`:

```
1 $ ls random/ | head -n 5
```


stdin

The opposite of `stdout` is `stdin`: standard input. Some programs read from `stdin` when they expect a path but aren't given any.

Try this in 03-pipes/:

```
1 $ ls random/ | head -n 5
```

Observation

Normally `head` expects a filename, but when none is given, it falls back to `stdin` — which is what `ls` printed to `stdout`.

Convention

We sometimes call the vertical bar (`|`) the **pipe** character.

The power of pipes

You can chain commands with pipes. Classic recipe (still in 03-pipes):

```
1 $ cat numbers
2 $ cat numbers | sort
3 $ cat numbers | sort | uniq
4 $ cat numbers | sort | uniq | wc
```

The power of pipes

You can chain commands with pipes. Classic recipe (still in 03-pipes):

```
1 $ cat numbers
2 $ cat numbers | sort
3 $ cat numbers | sort | uniq
4 $ cat numbers | sort | uniq | wc
```

Observations

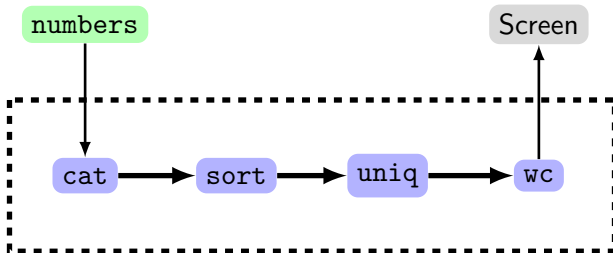
- ▶ Each program takes the last one's stdout as stdin
- ▶ Only the final program will print to terminal

Explanation

- ▶ By default, sort sorts stdin in dictionary order (check man page for more)
- ▶ wc is short for “word count”, although the first thing it prints is the number of lines.

Pipes illustrated

```
1 $ cat numbers | sort | uniq | wc
```



Challenge

Can you think of a way to eliminate a pipe? (Hint: `man sort`)

grep

Try this in 03-pipes: ⁸

```
1 $ ls random/ | grep JI
```

⁸For Mac users, instead of `grep` you might have to type `ggrep`

grep

Try this in 03-pipes: ⁸

```
1 $ ls random/ | grep JI
```

Explanation

grep is a powerful tool to match a substring. By default, it takes a file (or stdin), and prints all lines containing a pattern ("JI").

⁸For Mac users, instead of grep you might have to type ggrep

Challenge

Inside 03-pipes/random/:

- ▶ List all filenames containing "FDU"
- ▶ List all filenames containing "UM" (upper and lower cases)
(Hint: man grep)
- ▶ List all filenames containing "UM" but not "FDU" (upper and lower cases for both substrings)

Solution

```
1 $ ls | grep FDU
2 $ ls | grep -i UM
3 $ ls | grep -i UM | grep -i -v FDU
```


Conclusion

- ▶ Files and directories form a tree
- ▶ One tool does one thing, but flags specify how
- ▶ When in doubt, read documentation
- ▶ Chain together tools and unleash immense power

The End

Thank You For Coming!

Credits

- ▶ The Free Software Foundation, logo of GNU Bash.
<https://commons.wikimedia.org/wiki/File:Gnu-bash-logo.svg>
- ▶ Robin Nicholas, Patricia Saunders and The Open Group, logo of UNIX.
https://commons.wikimedia.org/wiki/File:UNIX_logo.svg
- ▶ The Free Software Foundation, logo of GNU.
https://commons.wikimedia.org/wiki/File:The_GNU_logo.png
- ▶ Larry Ewing, logo of Linux.
https://commons.wikimedia.org/wiki/File:Linux_logo.jpg
- ▶ Poul-Henning Kamp, Beastie.
<https://commons.wikimedia.org/wiki/File:Daemon-phk.svg>
- ▶ ZxxZxxZ. Linux command-line. https://commons.wikimedia.org/wiki/File:Linux_command-line._Bash._GNOME_Terminal._screenshot.png